

# Multiple Linear Regression Examples

*Cody Frisby*

*February 12, 2016*

Supervisor performance example shared in class: Companies are continually evaluating supervisors to, not only determine adequate performance, but also gauge employee morale (an important indicator for employee productivity). In an effort to understand the important aspects of a satisfactory supervisor, employees at a certain company were asked to provide an overall rating and scores on 6 characteristics of their immediate managers. Namely, employees were asked to rate the following statements on a scale from 0 to 100.

Variable	Description
Rating	Overall rating of the supervisor performance
Complaints	Score for “Your supervisor handles employee complaints appropriately.”
Privileges	Score for “Your supervisor allows special privileges.”
Learn	Score for “Your supervisor provides opportunities to learn new things.”
Raises	Score for “Your supervisor bases raises on performance.”
Critical	Score for “Your supervisor is too critical of poor performance.”
Advance	Score for “I am not satisfied with the rate I am advancing in the company.”

If we are going to get started doing multiple linear regression we are going to need to review some linear algebra concepts. Chapter 1 section 8 in the Graybill book reviews the concepts. Review matrices and vector operations.

**A note for doing linear algebra operations in R.** There is an excellent library that mimics a lot of the same syntax as matlab. Install the `pracma` library by typing in `install.packages(“pracma”)` and after that finishes installing call the library by `library(pracma)`.

Rendering matrices using knitr and R markdown.

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

There is an example on page 225 (4.2.1) in the Graybill book. The `PLASTICS.DAT` file in the `datafiles` directory contains the data.

```
plastic <- read.table("~/Documents/MATH3710/datafiles/PLASTIC.DAT", header = FALSE)
# we need the column names added.
plastic <- plastic[,2:4] # drops the column not needed.
colnames(plastic) <- c("strength", "temp", "pressure")

# snap shot of the data
(head(plastic))
```

```
## strength temp pressure
## 1      30.7  240      16
## 2      24.7  250      18
## 3      30.6  260      16
## 4      32.8  240      10
## 5      20.7  240      20
## 6      34.5  260      16
```

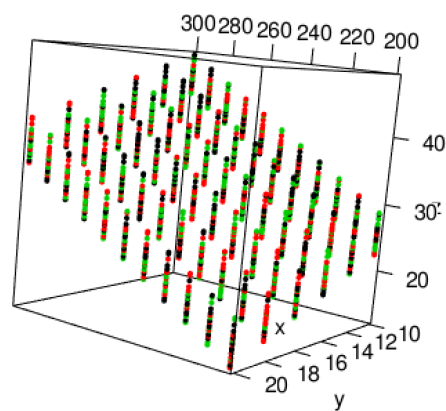
```
fit.plastic <- lm(strength ~ temp + pressure, data = plastic)
# load this library, run install.packages("dplyr") if you do not have it.
library(dplyr)
plastic.summary <- plastic %>%
  group_by(temp, pressure) %>%
  summarise_each(funs(mean, sd), strength)
plastic.summary
```

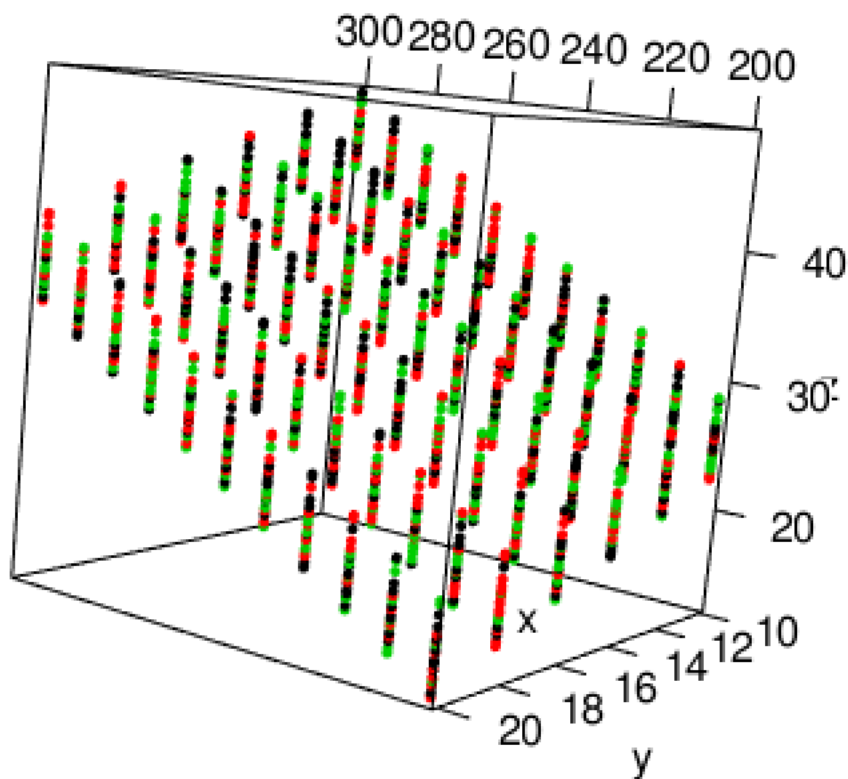
```
## Source: local data frame [66 x 4]
## Groups: temp [?]
##
##   temp pressure  mean      sd
##   (int)   (int) (dbl)   (dbl)
## 1    200      10    25 1.742843
## 2    200      12    23 1.742843
## 3    200      14    21 1.742843
## 4    200      16    19 1.742843
## 5    200      18    17 1.742843
## 6    200      20    15 1.742843
## 7    210      10    27 1.742843
## 8    210      12    25 1.742843
## 9    210      14    23 1.742843
## 10   210      16    21 1.742843
## .. ...      ...    ...      ...
```

Note, the difference between the text's standard deviation and this one is due to a sample and population calculation. R default is to calculate the sample standard deviation.

Playing around with R 3d plotting library:

Image of a 3D plot of the three variables:





Back to the Supervisor data. Let's get it loaded into R first, and repeat some of the steps shared in the SAS code from class:

```

sup <- read.table("~/Documents/MATH3710/multLR/Supervisor.txt", header = TRUE)
sup.pre <- sup
# SAS code :
# w = Complaints + Learn;
# v = Complaints - Learn;
# yprime = Ratings - Learn;
sup$w <- sup$Complaints + sup$Learn
sup$v <- sup$Complaints - sup$Learn
sup$yprime <- sup$Rating - sup$Learn
# SAS:
# proc reg data = supervisor;
# model yprime = v;
fit1 <- lm(Rating ~ Complaints+Privileges+Critical, data = sup)
fit2 <- lm(Rating ~ Learn+Raises+Advance, data = sup)
fit3 <- lm(Rating ~ w, data = sup)
fit4 <- lm(yprime ~ v, data = sup)
fit5 <- lm(Rating ~ Complaints+Learn, data=sup)
# here is the model with all variables before adding v and w
fit6 <- lm(Rating ~ ., data = sup.pre)

```

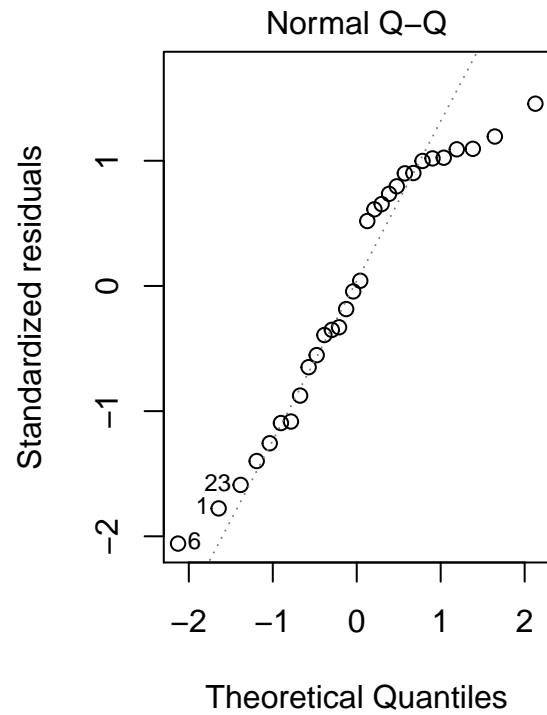
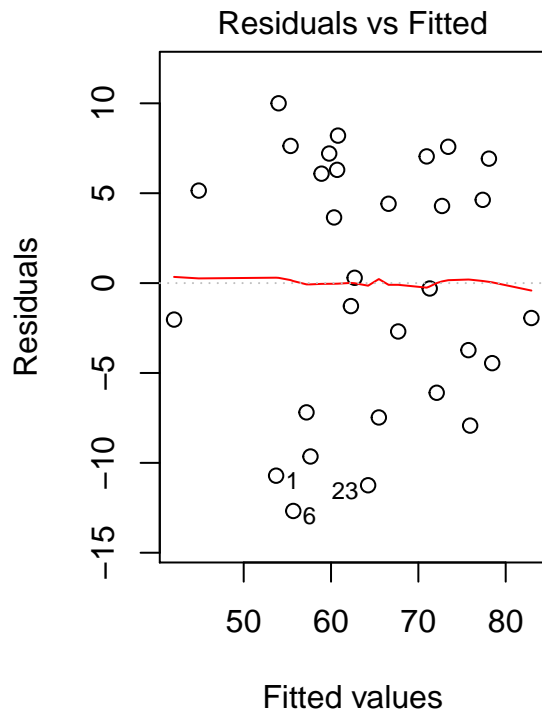
## Diagnostics of all the above models:

### Fit model 1

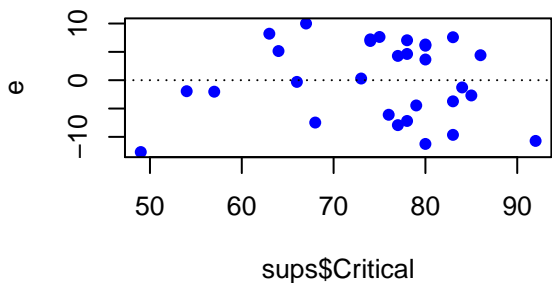
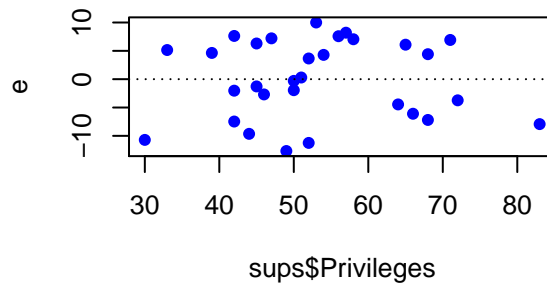
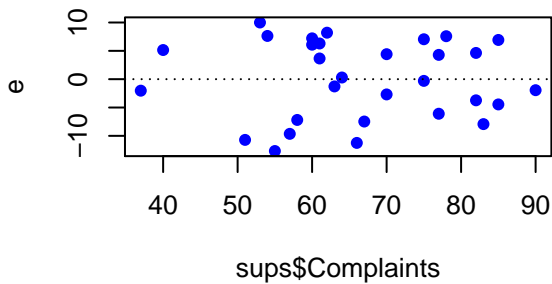
```
# summary(fit1)$r.squared # just print r squared value
# summary(fit1)$coeff # just print the model coefficients
summary(fit1) # prints summary of model

##
## Call:
## lm(formula = Rating ~ Complaints + Privileges + Critical, data = sups)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -12.676  -5.689  -0.003   6.246  10.002
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 15.027624  11.535858   1.303   0.204
## Complaints   0.779814   0.122678   6.357 9.88e-07 ***
## Privileges  -0.050392   0.132571  -0.380   0.707
## Critical     0.004649   0.138465   0.034   0.973
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.237 on 26 degrees of freedom
## Multiple R-squared:  0.6831, Adjusted R-squared:  0.6465
## F-statistic: 18.68 on 3 and 26 DF,  p-value: 1.144e-06

par(mfrow = c(1,2)) # sets up side by side plots
plot(fit1, which = c(1,2))
```



```
e <- residuals(fit1)
# plot residuals by all predictors
par(mfrow = c(2,2))
plot(sups$Complaints, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Privileges, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Critical, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```

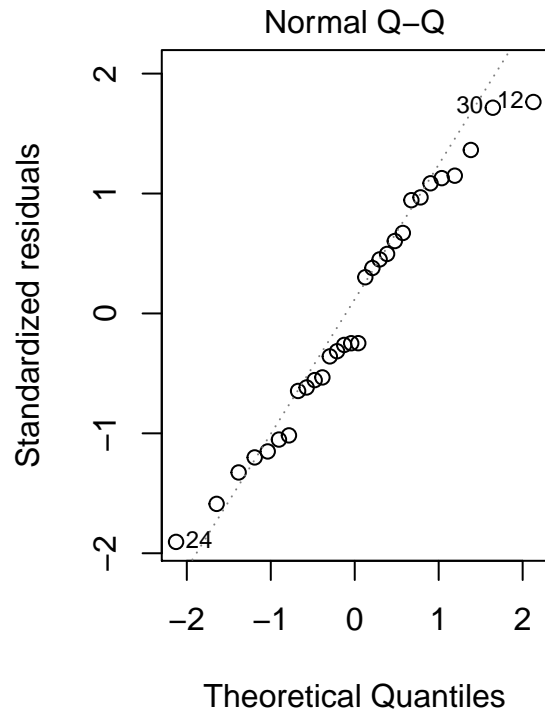
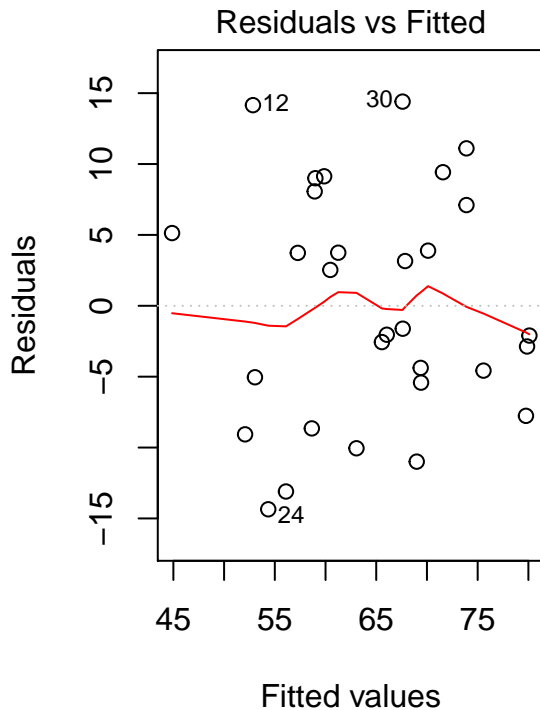


## Fit model 2

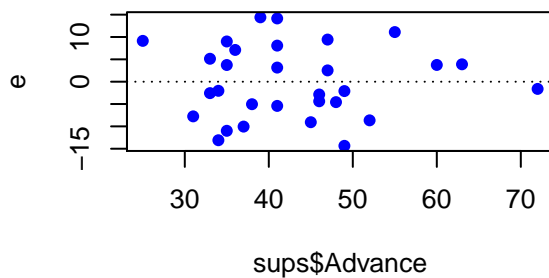
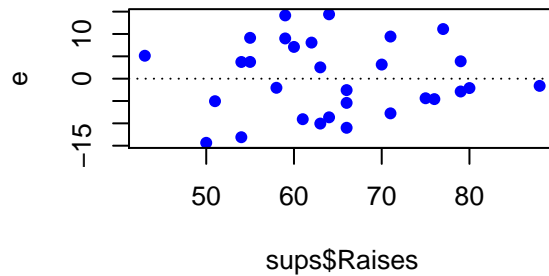
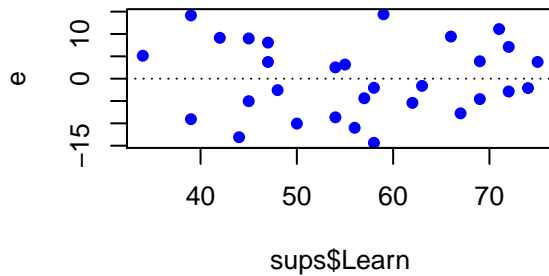
```
summary(fit2) # prints summary of model
```

```
##
## Call:
## lm(formula = Rating ~ Learn + Raises + Advance, data = sups)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.354  -5.324  -1.828   6.612  14.403
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  17.6511    10.2245   1.726  0.09615 .
## Learn         0.5484     0.1835   2.988  0.00606 **
## Raises       0.5658     0.2143   2.640  0.01383 *
## Advance      -0.4774     0.1964  -2.431  0.02228 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.601 on 26 degrees of freedom
## Multiple R-squared:  0.5524, Adjusted R-squared:  0.5007
## F-statistic: 10.69 on 3 and 26 DF,  p-value: 9.262e-05
```

```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit2, which = c(1,2))
```



```
e <- residuals(fit2)
# plot residuals by all predictors
par(mfrow = c(2,2))
plot(sups$Learn, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Raises, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Advance, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```



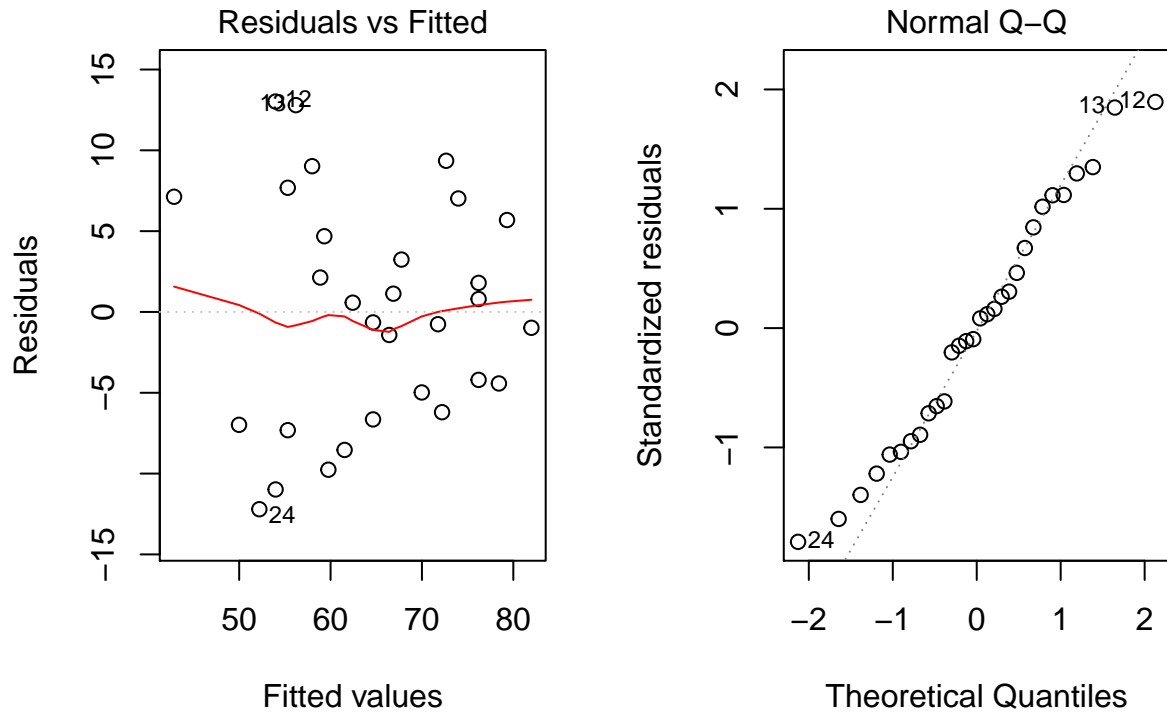
### Fit model 3

```
summary(fit3) # prints summary of model
```

```
##
## Call:
## lm(formula = Rating ~ w, data = sups)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -12.2052  -5.8973  -0.0372   5.4364  13.0172
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.98821    7.38841   1.352   0.187
## w             0.44439    0.05914   7.514 3.49e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.133 on 28 degrees of freedom
## Multiple R-squared:  0.6685, Adjusted R-squared:  0.6566
```

```
## F-statistic: 56.46 on 1 and 28 DF,  p-value: 3.487e-08
```

```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit3, which = c(1,2))
```

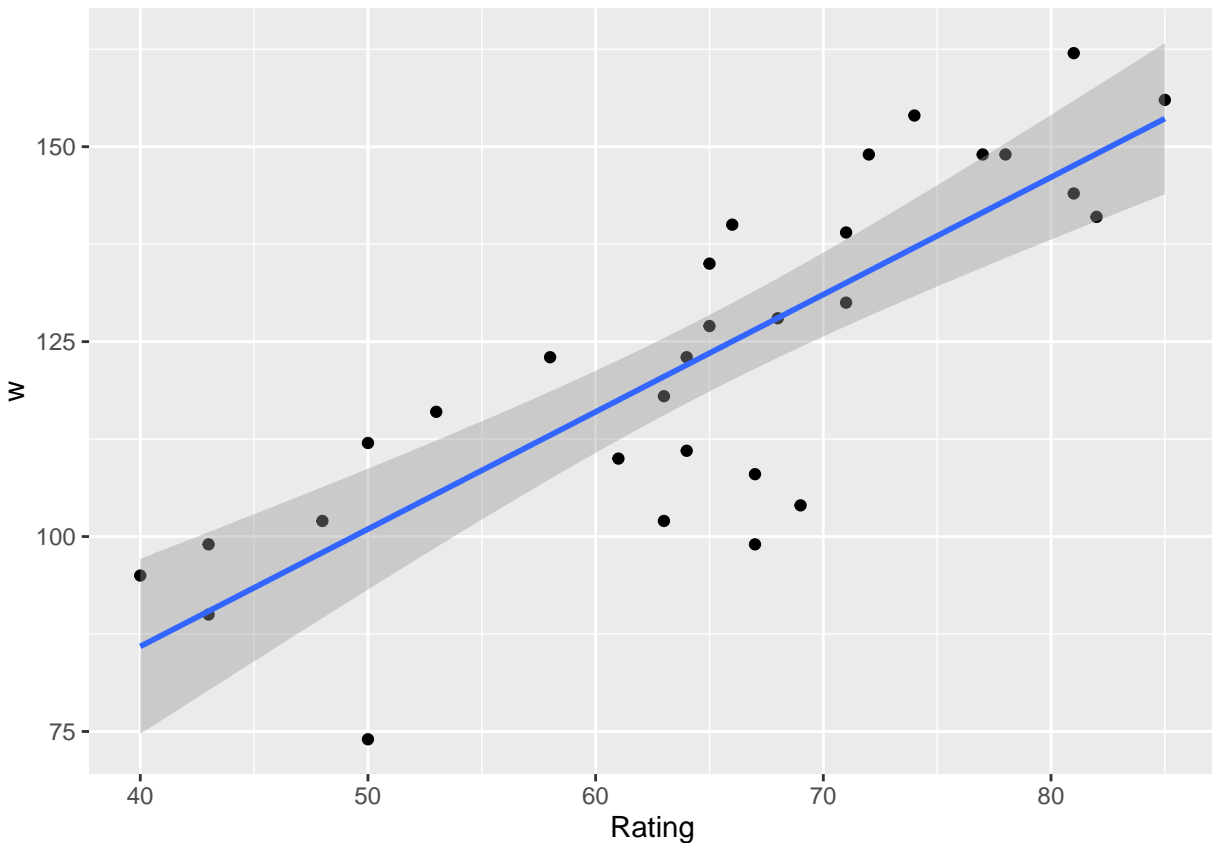


```
e <- residuals(fit3)
# plot residuals by all predictors
# redundant plot below as above
#plot(sups$w, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```

```
library(ggplot2)
g <- ggplot(sups, aes(x = Rating, y = w))
g <- g + geom_point()
g <- g + stat_smooth(method = "lm", formula = y ~ x)
g
```

Fit Plot for model 3





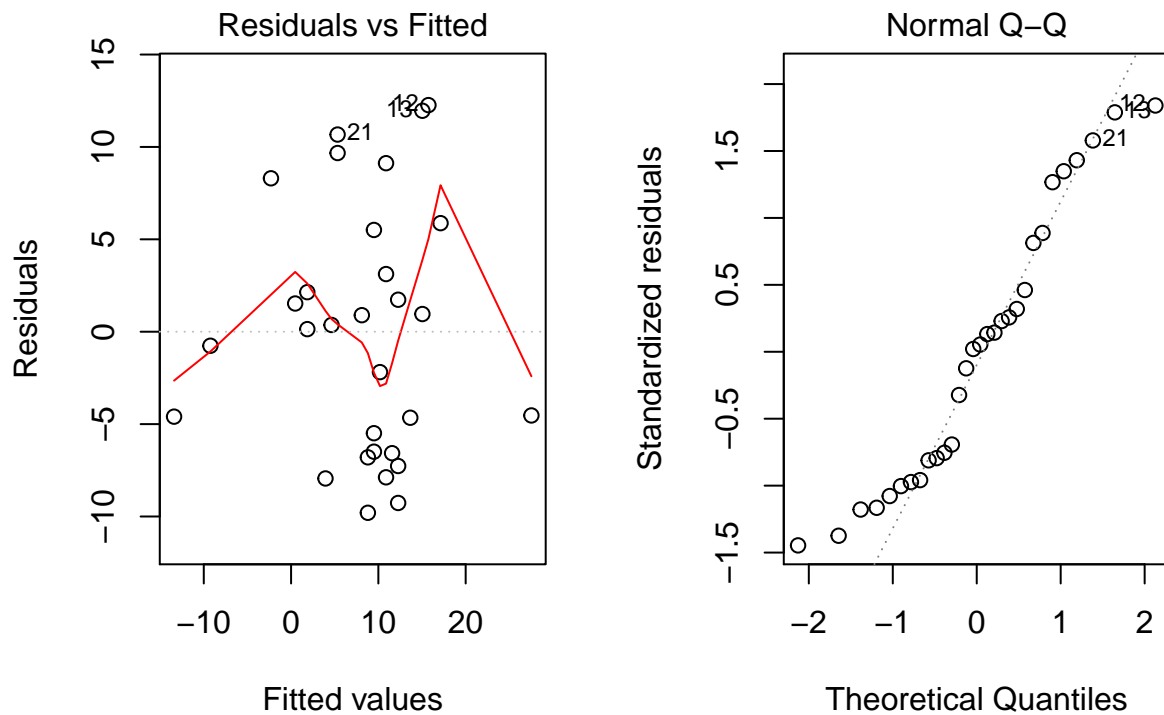
*Note, ggplot automatically adds confidence bands to the plot.*

#### Fit model 4

```
summary(fit4) # prints summary of model
```

```
##
## Call:
## lm(formula = yprime ~ v, data = sups)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.799 -6.242  0.252  4.911 12.263
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.1665     1.7079   0.683    0.5
## v             0.6938     0.1129   6.147 1.23e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.891 on 28 degrees of freedom
## Multiple R-squared:  0.5744, Adjusted R-squared:  0.5592
## F-statistic: 37.79 on 1 and 28 DF, p-value: 1.233e-06
```

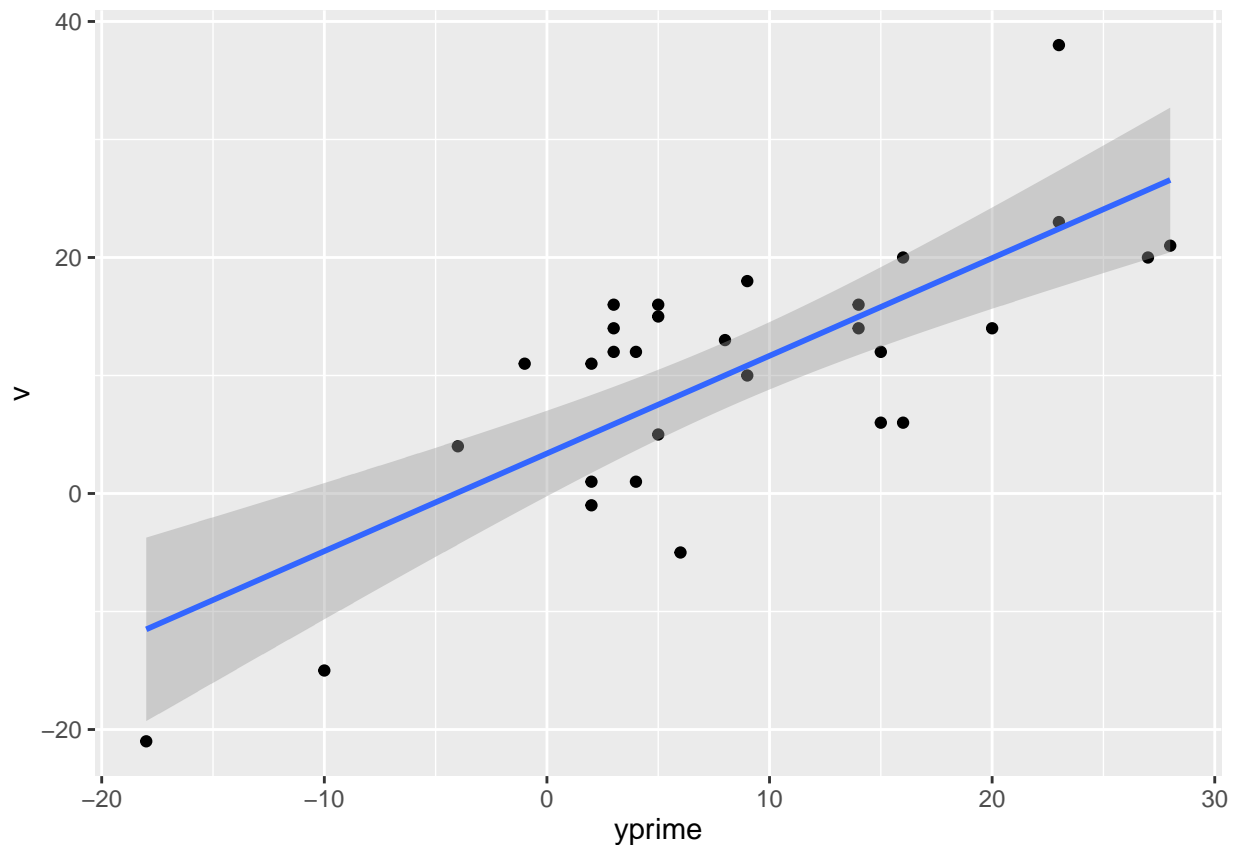
```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit4, which = c(1,2))
```



```
e <- residuals(fit4)
# plot residuals by all predictors
# redundant plot below as above
#plot(sups$w, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```

```
library(ggplot2)
g <- ggplot(sups, aes(x = yprime, y = v))
g <- g + geom_point()
g <- g + stat_smooth(method = "lm", formula = y ~ x)
g
```

Fit Plot for model 4

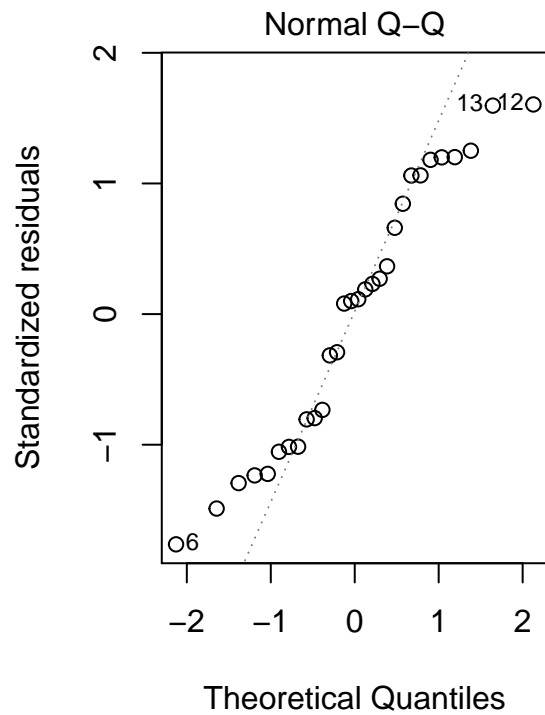
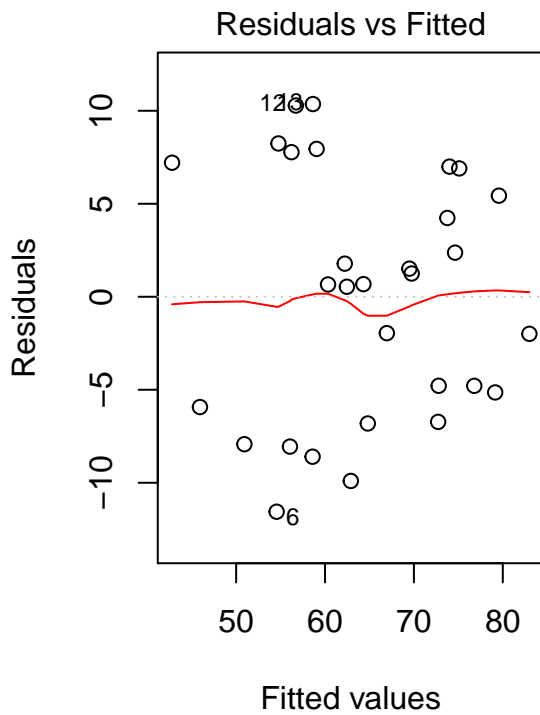


## Fit model 5

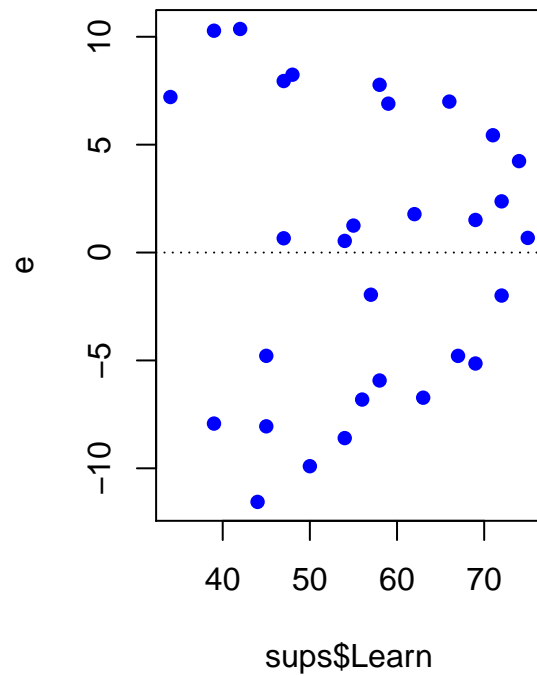
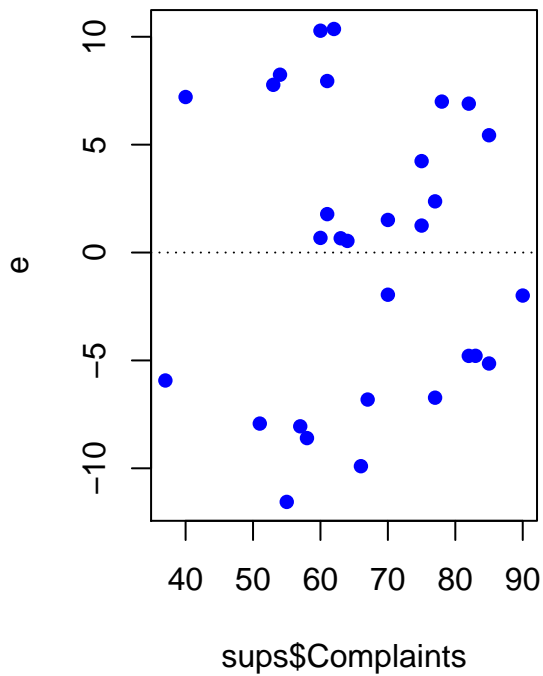
```
summary(fit5) # prints summary of model
```

```
##
## Call:
## lm(formula = Rating ~ Complaints + Learn, data = sups)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.5568  -5.7331   0.6701   6.5341  10.3610
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.8709     7.0612   1.398   0.174
## Complaints     0.6435     0.1185   5.432 9.57e-06 ***
## Learn          0.2112     0.1344   1.571   0.128
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.817 on 27 degrees of freedom
## Multiple R-squared:  0.708, Adjusted R-squared:  0.6864
## F-statistic: 32.74 on 2 and 27 DF, p-value: 6.058e-08
```

```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit5, which = c(1,2))
```



```
e <- residuals(fit5)
# plot residuals by all predictors
par(mfrow = c(1,2))
plot(sups$Complaints, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Learn, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```

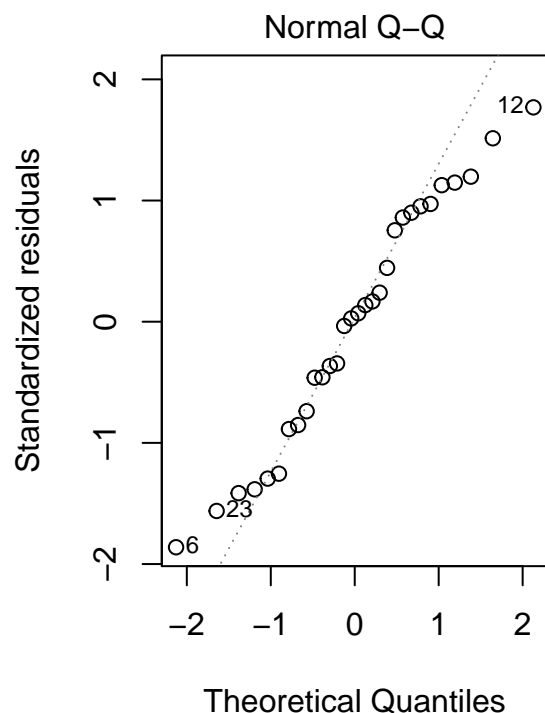
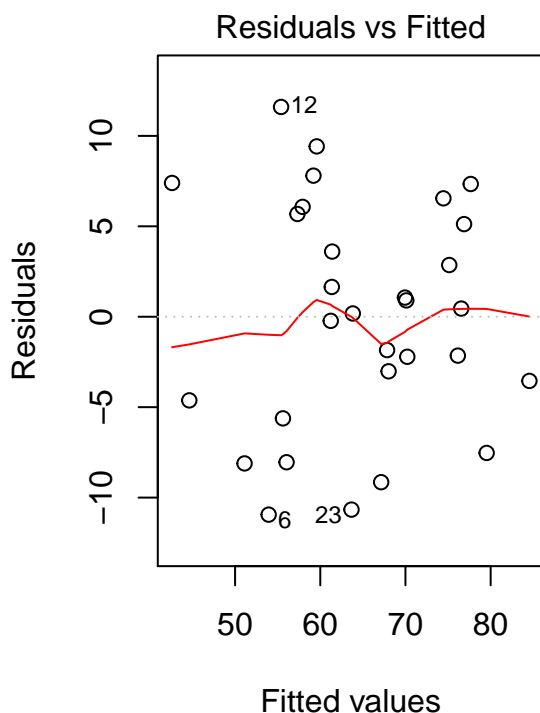


## Fit model 6

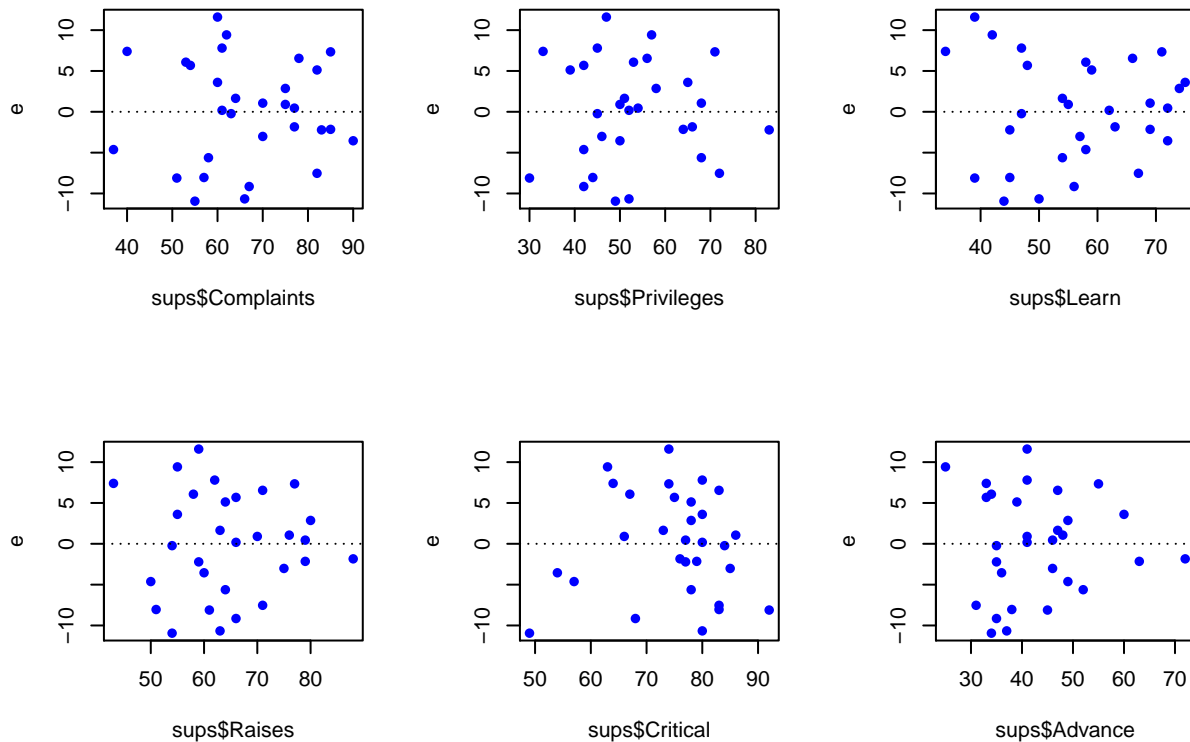
```
summary(fit6) # prints summary of model
```

```
##
## Call:
## lm(formula = Rating ~ ., data = sups.pre)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.9418  -4.3555   0.3158   5.5425  11.5990
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  10.78708    11.58926   0.931  0.361634
## Complaints    0.61319     0.16098   3.809  0.000903 ***
## Privileges   -0.07305     0.13572  -0.538  0.595594
## Learn         0.32033     0.16852   1.901  0.069925 .
## Raises       0.08173     0.22148   0.369  0.715480
## Critical      0.03838     0.14700   0.261  0.796334
## Advance      -0.21706     0.17821  -1.218  0.235577
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.068 on 23 degrees of freedom
## Multiple R-squared:  0.7326, Adjusted R-squared:  0.6628
## F-statistic: 10.5 on 6 and 23 DF,  p-value: 1.24e-05
```

```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit6, which = c(1,2))
```



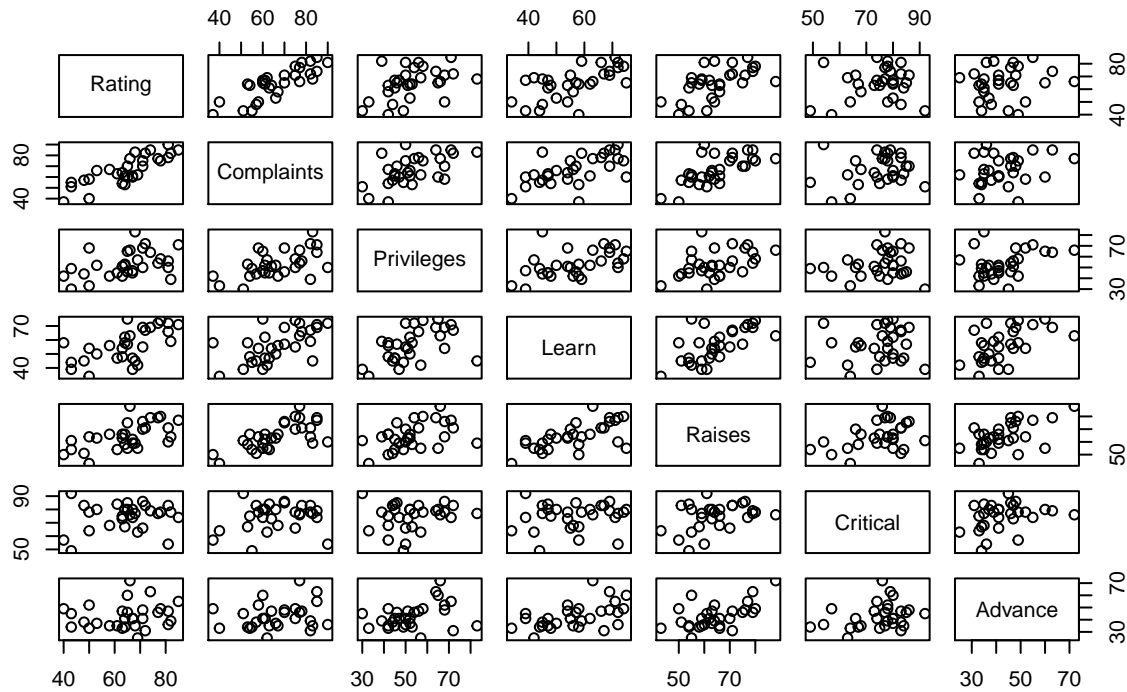
```
e <- residuals(fit6)
# plot residuals by all predictors
par(mfrow = c(2,3))
plot(sups$Complaints, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Privileges, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Learn, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Raises, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Critical, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Advance, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```



Here's a scatterplot matrix of the supervisor data without the added variables

```
# note, I created this variable before adding all the additional vars.
plot(sups.pre, main = "Scatterplot Matrix")
```

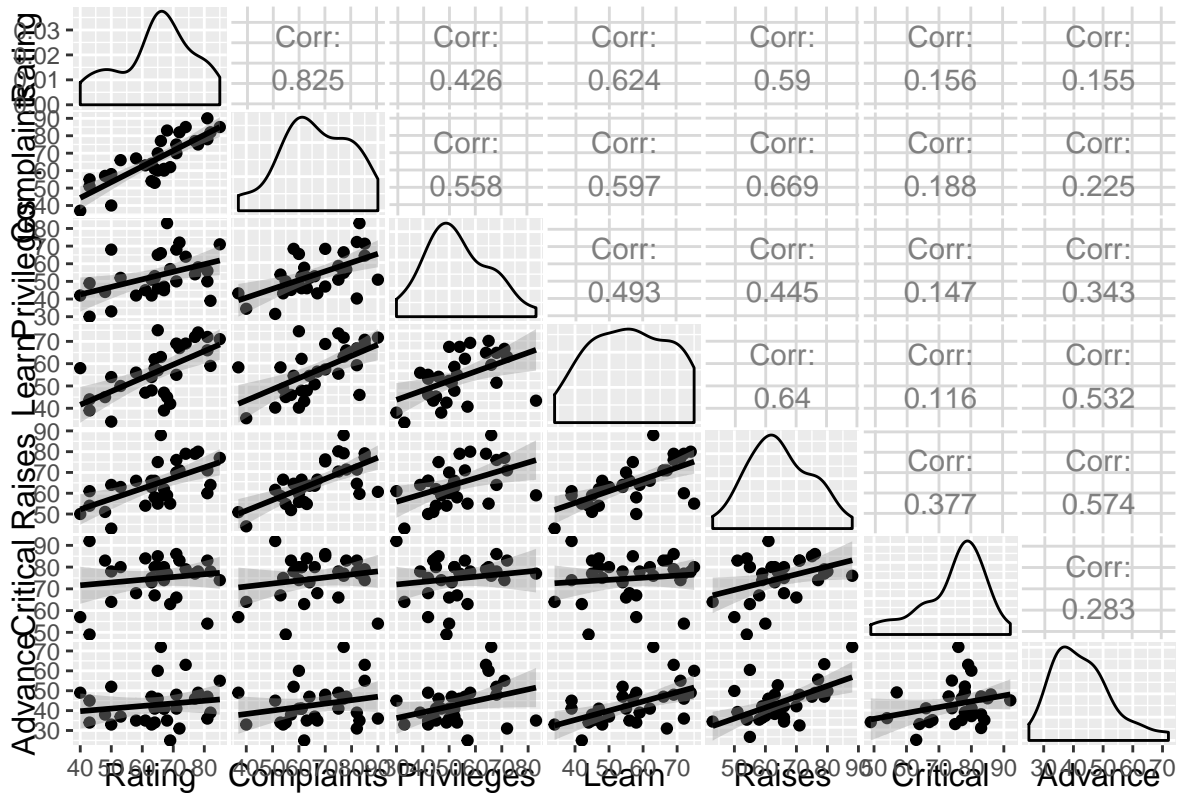
## Scatterplot Matrix



Notice the multi-collinearity of a few of the x variables with each other. Particularly Learn vs Raises.

Here, I'd recommend installing the GGally package. It is a companion package to ggplot2. This next plot is slow but very informative for multi linear regression.

```
library(GGally)
gsups <- ggpairs(sups.pre, lower = list(continuous = "smooth"))
gsups
```



Correlation coefficients for our supervisor data. Note that this is a part of the above plot. I do this below just to show you how its done in R.

```
cor(sups.pre)
```

```
##           Rating Complaints Privileges   Learn   Raises  Critical
## Rating      1.000000  0.8254176  0.4261169 0.6236782 0.5901390 0.1564392
## Complaints  0.8254176  1.0000000  0.5582882 0.5967358 0.6691975 0.1877143
## Privileges  0.4261169  0.5582882  1.0000000 0.4933310 0.4454779 0.1472331
## Learn       0.6236782  0.5967358  0.4933310 1.0000000 0.6403144 0.1159652
## Raises      0.5901390  0.6691975  0.4454779 0.6403144 1.0000000 0.3768830
## Critical    0.1564392  0.1877143  0.1472331 0.1159652 0.3768830 1.0000000
## Advance     0.1550863  0.2245796  0.3432934 0.5316198 0.5741862 0.2833432
##           Advance
## Rating      0.1550863
## Complaints  0.2245796
## Privileges  0.3432934
## Learn       0.5316198
## Raises      0.5741862
## Critical    0.2833432
## Advance     1.0000000
```

Now, using the model to make predictions:

```
# test the glm model
test <- data.frame(Complaints = c(61,71,61), Privileges = c(45,45,45),
  Learn = c(56,56,56), Raises = c(71,71,81), Critical = c(57,57,57), Advance = c(25,25,25))
```



```
fit.all <- glm(Rating ~ ., data = sups.pre)
sup.pred <- predict.glm(fit.all, newdata = test, interval =
  "confidence", se.fit = TRUE)
preds <- cbind(sup.pred$fit, sup.pred$se.fit)
colnames(preds) <- c("estimate", "std error")
preds
```

```
## estimate std error
## 1 65.40717 5.257481
## 2 71.53905 4.712535
## 3 66.22449 7.059975
```

```
summary(fit.all)$coeff # print the coefficients from the glm model.
```

```
##           Estimate Std. Error   t value   Pr(>|t|)
## (Intercept) 10.78707639 11.5892572  0.9307824 0.3616337210
## Complaints   0.61318761  0.1609831  3.8090182 0.0009028679
## Privileges  -0.07305014  0.1357247 -0.5382229 0.5955939205
## Learn        0.32033212  0.1685203  1.9008516 0.0699253459
## Raises       0.08173213  0.2214777  0.3690310 0.7154800884
## Critical     0.03838145  0.1469954  0.2611064 0.7963342642
## Advance     -0.21705668  0.1782095 -1.2179862 0.2355770486
```

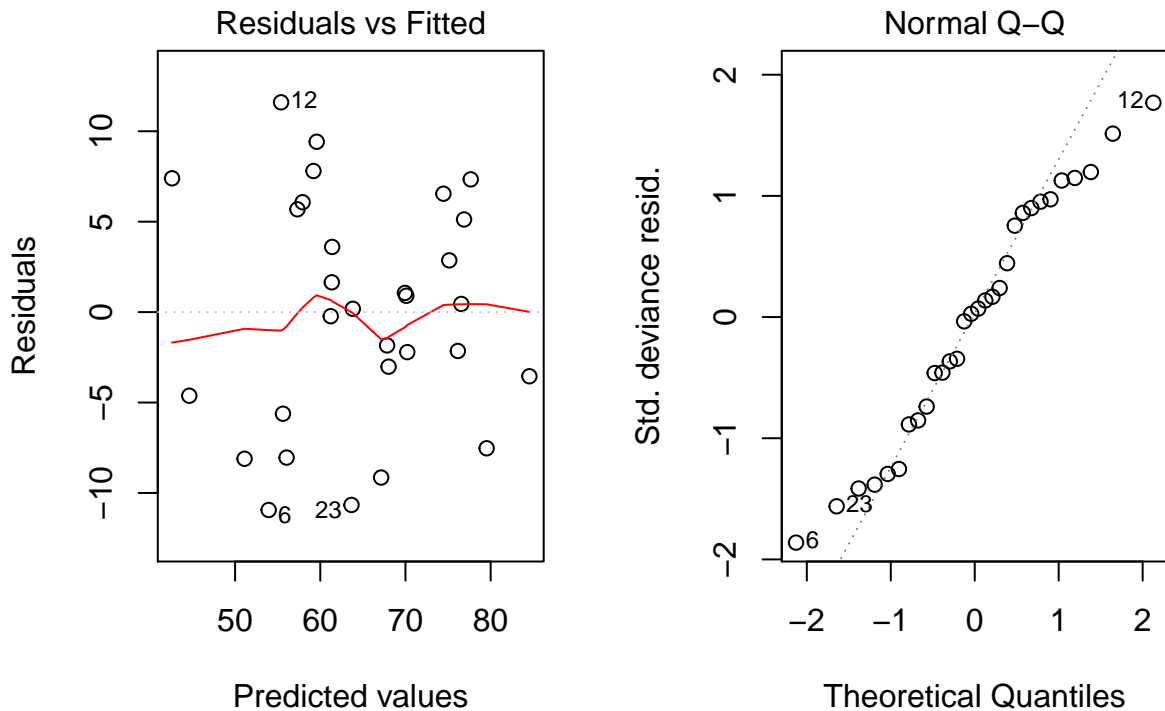
Summary of the glm model from above

```
summary(fit.all) # prints summary of model
```

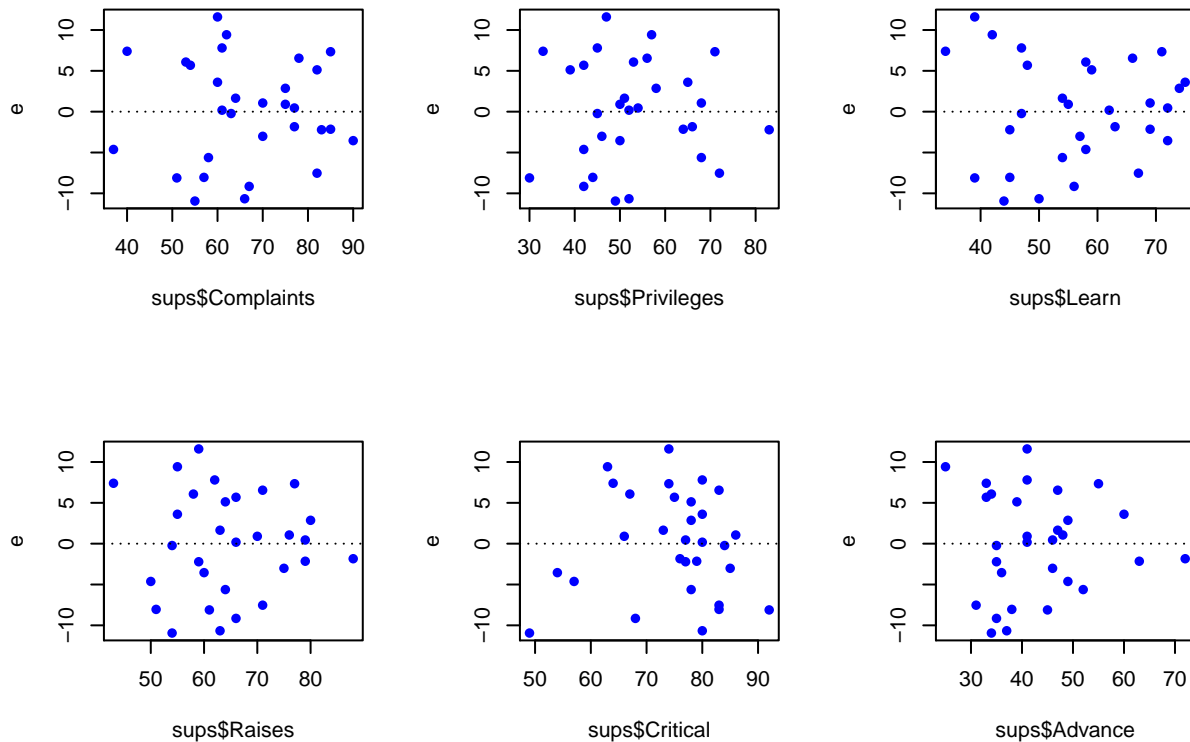
```
##
## Call:
## glm(formula = Rating ~ ., data = sups.pre)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -10.9418  -4.3555   0.3158   5.5425  11.5990
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.78708   11.58926   0.931 0.361634
## Complaints   0.61319   0.16098   3.809 0.000903 ***
## Privileges  -0.07305   0.13572  -0.538 0.595594
## Learn        0.32033   0.16852   1.901 0.069925 .
## Raises       0.08173   0.22148   0.369 0.715480
## Critical     0.03838   0.14700   0.261 0.796334
## Advance     -0.21706   0.17821  -1.218 0.235577
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 49.95654)
##
##      Null deviance: 4297  on 29  degrees of freedom
## Residual deviance: 1149  on 23  degrees of freedom
```

```
## AIC: 210.5
##
## Number of Fisher Scoring iterations: 2
```

```
par(mfrow = c(1,2)) # sets up side by side plots
plot(fit.all, which = c(1,2))
```



```
e <- residuals(fit.all)
# plot residuals by all predictors
par(mfrow = c(2,3))
plot(sups$Complaints, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Privileges, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Learn, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Raises, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Critical, e, col = "blue", pch = 16); abline(h=0, lty = 3)
plot(sups$Advance, e, col = "blue", pch = 16); abline(h=0, lty = 3)
```



Now we need to add the 10,000 iterations like professor did using SAS. I don't yet know how to do this without a for loop in R so this is a slow block of code.

```
bias <- vector()
sups$index <- 1:30
for (i in 1:10000) {
  ind <- sample(1:dim(sups)[1], 10)
  # now we can subset the carmpg data frame into 2 data frames using this index.
  test <- sups[sups$index %in% ind, ]
  training <- sups[!(sups$index %in% ind), ]
  # now to fit a new model with our subsetting training data.
  fit.train <- lm(Rating ~
    Complaints+Privileges+Learn+Raises+Critical+Advance,
    data = training)
  testing <- predict.lm(fit.train, newdata = test)
  e <- test$Rating - testing
  e.squared <- e^2
  p.i <- c(mean(e), mean(e.squared), sqrt(mean(e.squared)))
  bias <- rbind(bias, p.i)
}
colnames(bias) <- c("pbias", "pmse", "rpmse")
#mean(bias)
apply(bias, 2, mean) # takes the means of the columns of our matrix
```

```
##      pbias      pmse      rpmse
## -0.5666259 78.3816421 8.6821439
```

```
hist(bias[,1], col = "green", breaks = 20, main = "Histogram of Model Bias", xlab = "observed - predicted")
```

