

# Computer Repairs

*ccfrisby*

*January 4, 2016*

Let's get started with simple linear regression using R. This first data set is small so we have the data loaded into two variables: minutes and units. We then put them together with the cbind function and store the result in a new variable called repair.

```
units <- c(1,2,3,4,4,5,6,6,7,8,9,9,10,10,3)
minutes <- c(23,29,49,64,74,87,96,97,109,119,149,145,154,166, NA)
repair <- as.data.frame(cbind(units, minutes))
repair <- repair[complete.cases(repair), ]
# note, the above line removes missing values.
describe(repair)
```


```
##          vars  n  mean    sd median trimmed   mad min max range  skew
## units         1 14  6.00  2.96   6.0   6.08  3.71   1  10     9 -0.12
## minutes        2 14 97.21 46.22  96.5  97.67 59.30  23 166   143 -0.09
##          kurtosis    se
## units        -1.43  0.79
## minutes       -1.37 12.35
```

Above is a summary of the data table “repair”. Below is a table of the correlation coefficients.

```
cor(repair, method = "pearson", use = "complete.obs")
```

```
##          units  minutes
## units  1.0000000 0.9936987
## minutes 0.9936987 1.0000000
```

Here's a scatter plot of the data.



Now to display a summary of the model. Above we fit a model with minutes being function of units. We stored this model in a variable called m.

```
m <- lm(minutes ~ units, data = repair)
anova(m) #ANOVA table
```

```
## Analysis of Variance Table
##
## Response: minutes
##      Df Sum Sq Mean Sq F value    Pr(>F)
## units     1 27419.5  27419.5   943.2 8.916e-13 ***
## Residuals 12   348.8    29.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(m) #summary of linear model
```

```
##
## Call:
## lm(formula = minutes ~ units, data = repair)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.2318 -3.3415 -0.7143  4.7769  7.8033
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.162      3.355     1.24   0.239
## units         15.509      0.505    30.71 8.92e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



```
##
## Residual standard error: 5.392 on 12 degrees of freedom
## Multiple R-squared:  0.9874, Adjusted R-squared:  0.9864
## F-statistic: 943.2 on 1 and 12 DF,  p-value: 8.916e-13
```

Here we attempt to reproduce some of the output from the Proc Reg command in SAS.

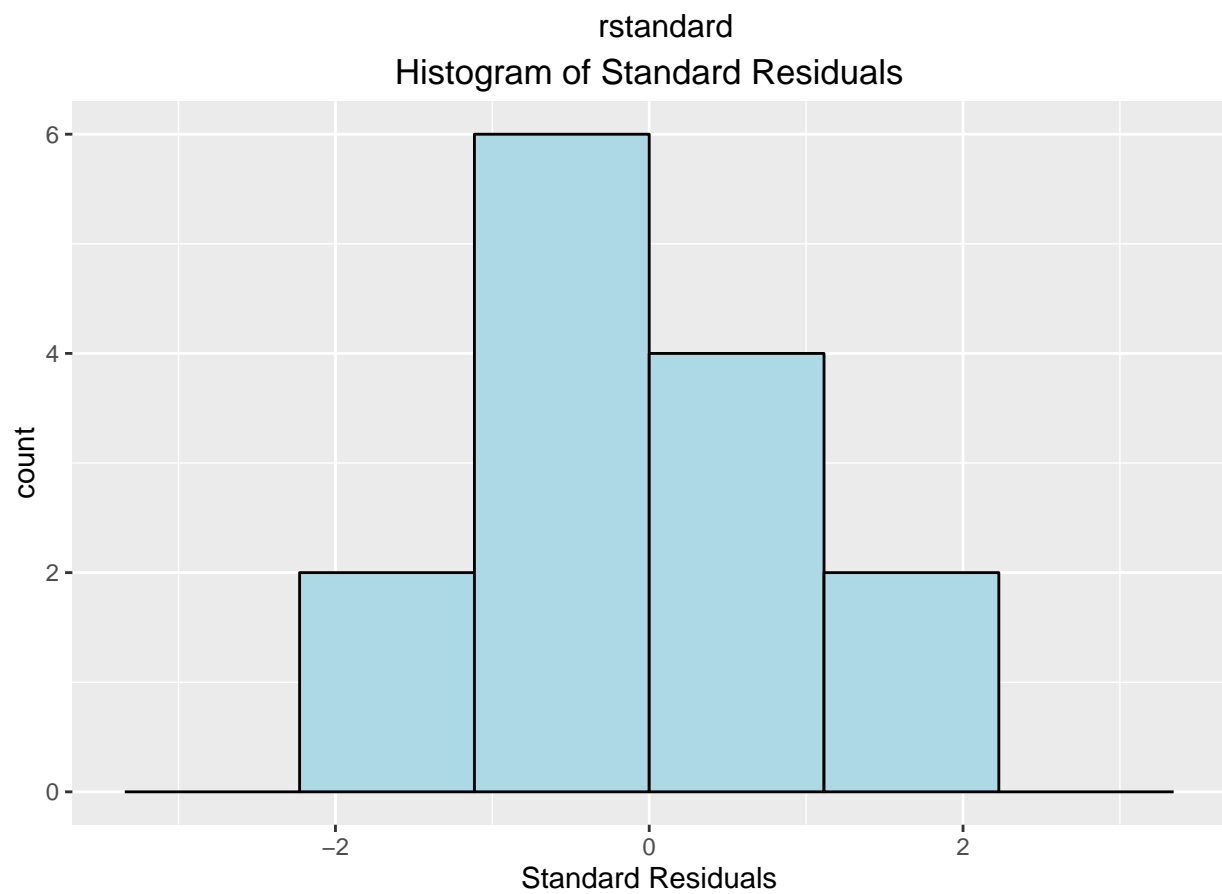
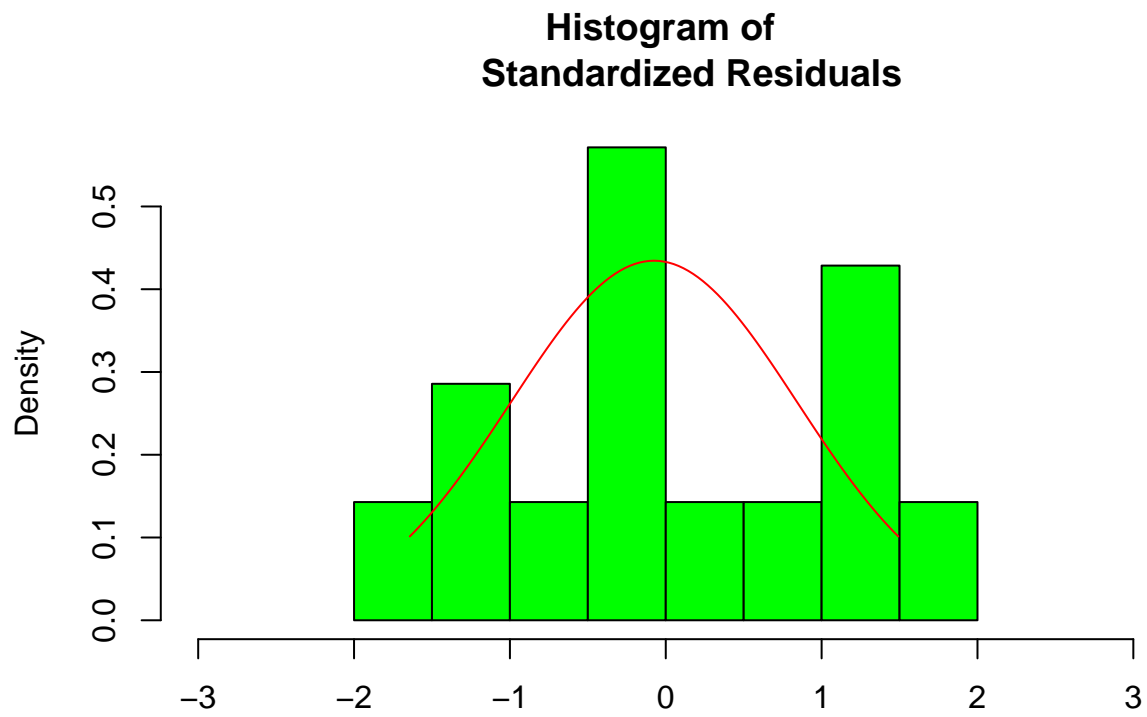
```
# add confidnece intervals
p <- predict.lm(m, se.fit = T, interval = "confidence")
psub <- p$fit[,2:3]
colnames(psub) <- c("95% Conf Lower", "95% Conf Upper")
s <- cbind(psub, stderror = p$se.fit)
# still need to add prediction intervals
# cbind is the command that binds columns together in either a
# data frame or matrix.
t <- cbind(repair$minutes, fitted.values(m), residuals(m))
colnames(t) <- c("observed", "predicted", "residuals")
cbind(t, s)
```

##	observed	predicted	residuals	95% Conf Lower	95% Conf Upper	stderror
## 1	23	19.67043	3.3295739	13.33625	26.00460	2.907169
## 2	29	35.17920	-6.1791980	29.77303	40.58537	2.481245
## 3	49	50.68797	-1.6879699	46.13246	55.24348	2.090821
## 4	64	66.19674	-2.1967419	62.36271	70.03077	1.759688
## 5	74	66.19674	7.8032581	62.36271	70.03077	1.759688
## 6	87	81.70551	5.2944862	78.37864	85.03239	1.526920
## 7	96	97.21429	-1.2142857	94.07462	100.35395	1.440999
## 8	97	97.21429	-0.2142857	94.07462	100.35395	1.440999
## 9	109	112.72306	-3.7230576	109.39618	116.04993	1.526920
## 10	119	128.23183	-9.2318296	124.39780	132.06586	1.759688
## 11	149	143.74060	5.2593985	139.18509	148.29611	2.090821
## 12	145	143.74060	1.2593985	139.18509	148.29611	2.090821
## 13	154	159.24937	-5.2493734	153.84321	164.65554	2.481245
## 14	166	159.24937	6.7506266	153.84321	164.65554	2.481245

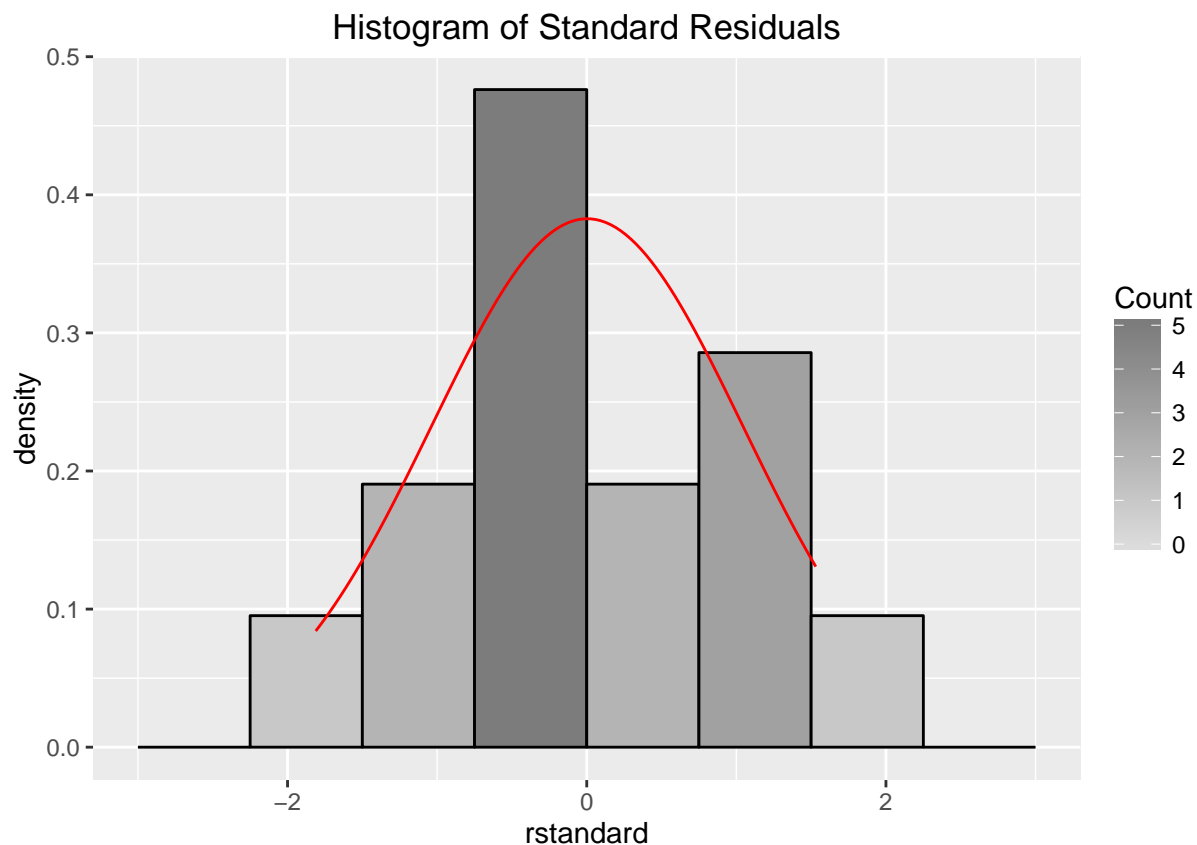
Below is a plot of the residuals vs fitted and the Normal Quantile plot.



Now to plot a histogram of the residuals from our model. We will do this three different ways for R illustration purposes.



Or we can plot the residuals this way, using ggplot.



Here we attempt to reproduce the same output as the Proc Univariate command in SAS. Here the variable of interest is from our model m, rstandard.

```
describe(rstandard(m)) #this is using the library psych
```

```
## vars n mean sd median trimmed mad min max range skew kurtosis
## 1 1 14 0 1.04 -0.14 0.03 1.36 -1.81 1.53 3.34 -0.06 -1.34
## se
## 1 0.28
```

```
t.test(rstandard(m))
```

```
##
## One Sample t-test
##
## data: rstandard(m)
## t = 0.0117, df = 13, p-value = 0.9908
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## -0.5986728 0.6051926
## sample estimates:
## mean of x
## 0.003259935
```

```
#or, using base R
x <- rstandard(m)
mean(x)
```

```
## [1] 0.003259935
```

```
sd(x)
```

```
## [1] 1.042519
```

```
summary(x)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -1.81100 -0.64770 -0.13750  0.00326  0.95120  1.53100
```

```
var(x)
```

```
## [1] 1.086846
```

```
u <- cbind(t, rstandard(m), hatvalues(m))
colnames(u) <- c("observed", "predicted", "residuals", "stdresid", "hatvals")
print(u)
```

```
##      observed predicted residuals   stdresid   hatvals
## 1         23  19.67043  3.3295739  0.73325354 0.29072682
## 2         29  35.17920 -6.1791980 -1.29086405 0.21177945
## 3         49  50.68797 -1.6879699 -0.33964371 0.15037594
## 4         64  66.19674 -2.1967419 -0.43103048 0.10651629
## 5         74  66.19674  7.8032581  1.53110484 0.10651629
## 6         87  81.70551  5.2944862  1.02388106 0.08020050
## 7         96  97.21429 -1.2142857 -0.23371442 0.07142857
## 8         97  97.21429 -0.2142857 -0.04124372 0.07142857
## 9        109 112.72306 -3.7230576 -0.71998832 0.08020050
## 10       119 128.23183 -9.2318296 -1.81140988 0.10651629
## 11       149 143.74060  5.2593985  1.05826626 0.15037594
## 12       145 143.74060  1.2593985  0.25340900 0.15037594
## 13       154 159.24937 -5.2493734 -1.09661925 0.21177945
## 14       166 159.24937  6.7506266  1.41023822 0.21177945
```

```
# confidence intervals and se values can be obtained in R by using
# predict.lm(m, se.fit = T, interval = "confidence")
# where m is your fitted model.
```